

# Embedded Systems Programming

Lecture 8 – Memory model and management

Jarno Tuominen

# Lecture 8 – Memory model and management

- Review
- Memory model of a C program
- Memory management

# Review of last lecture

- Functions (cont)
- The C preprocessor

# Review: Functions returning pointers

- C allows function to **return a pointer** to a
  - local variable (bad idea!) Why?
  - static variable
  - dynamically allocated memory
  - Function
- Syntax: add `"*"` in front of the function name to indicate that return value is a pointer (of type int, in the example)
- In this example the function returns an array – which is a pointer, remember?

```
*(rnds_p+0) : 1804289383
*(rnds_p+1) : 846930886
*(rnds_p+2) : 1681692777
*(rnds_p+3) : 1714636915
*(rnds_p+4) : 1957747793
```

```
int *getrandom(void) {
    static int r[5]; //must be static!
    for (int i = 0; i<5; i++) {
        r[i] = rand();
    }
    //Note: r is same as &r[0]
    return r;
}

int main() {
    int * rnds_p;
    rnds_p = getrandom();

    for (int i = 0; i<5; i++) {
        printf("*(%d) : %d\n", i, *(rnds_p+i));
    }
    return 0;
}
```

# Review: Function pointers

- A **function pointer** is a pointer variable that contains an **address of a function**, instead of a data object
- The syntax of declaration is similar to the syntax of declaring a function – but instead of using a function name, you use a pointer name inside the parenthesis
- Like with normal pointer variables, before using function pointer you need to assign it a value, i.e. the address of a function
- It is a good practice to type define declaration of function pointers as it will make your code much nicer
- Note: Function pointers are potentially very dangerous, as a loose pointer does not code access to wrong data, but if will cause your program to branch to a random address!

```
<return_type> (*<pointer_name>) (function_arguments);
```

```
typedef int (*fpComparer) (int x,int y);
```

```
int compare(int x,int y) {  
    ...  
}
```

Declare the function pointer and assign address of compare-function to it

```
int main() {  
    int result;  
    ...  
    fpComparer fpcomp = &compare;
```

```
result = fpcomp(a,b);  
//result = (*fpComparer)(a,b);  
...  
}
```

If not typedef'd

# Review: Function pointers packed in a struct

```
typedef uint8_t (*sensor_fp)(uint8_t SensorID, uint8_t param); ← Type define: sensor_fp  
//Struct for generic sensor instance  
typedef struct sensor_t {  
    char* name;  
    bool enabled;  
    sensor_fp init;  
    sensor_fp power_ctrl;  
} sensor_t;
```

Sensor-related data and its functions packed in a single "instance" (close to object-oriented thinking but still plain C!)

```
sensor_t sensors[MAX_SENSOR_COUNT]; Create an array of sensor instances  
sensors[0].name = "BME_1";  
sensors[0].enabled = true;  
sensors[0].init = bme160_initialize_sensor; } set-up the sensor instance  
sensors[0].power_ctrl = bme160_power_ctrl;  
sensors[1].name = "ECG";  
sensors[1].enabled = true;  
sensors[1].init = ads1293_init; } and another one  
sensors[1].power_ctrl = ads1293_power_ctrl;
```

Init all the sensors! Beautiful code <3

```
void init_sensors(uint8_t count) {  
    for (uint32_t i=0; i < count; i++)  
    {  
        e = sensors[i].init(i, NULL);  
    }  
}
```

# Review: functions with variable argument list

```
#include <stdio.h>
#include <stdarg.h>

double avg(int num,...) {
    va_list valist;
    double sum = 0.0;
    int i;
    /* initialize valist for num number of arguments */
    va_start(valist, num);
    /* access all the arguments assigned to valist */
    for (i = 0; i < num; i++) {
        sum += va_arg(valist, int);
    }
    va_end(valist); /* clean memory reserved for valist */
    return sum/num;
}

int main() {
    printf("Avg = %f\n", avg(4, 2,3,4,5));
    printf("Avg = %f\n", avg(3, 5,10,15));
}
```

num is number of arguments

va\_list is a data type that can hold list of arguments.  
Used by macros **va\_start**, **va\_arg** and **va\_end**

populate *valist* using macro **va\_start**

Get arguments of type **int**  
using macro **va\_arg**

**va\_end** releases memory

# Review: The C Preprocessor

- **C preprocessor** is the macro preprocessor, which provides the ability for the
  - inclusion of header files
  - macro expansions
  - conditional compilation
  - line control
  - Handling of pragma operators (in C99)
- invoked by the compiler as the first part of code translation
- Preprocessor macros begin with **#**

[https://en.wikipedia.org/wiki/C\\_preprocessor](https://en.wikipedia.org/wiki/C_preprocessor)

# Review: Preprocessor macros and directives

- **#include**
  - Inclusion of header files
  - Remember include guards
- **#define, #undef**
  - Defining/undefining macros
- **#if, #elif, #endif, #ifdef, #ifndef**
  - For conditional compilation
- **#error, #warning**
  - For custom errors/warnings

```
#ifndef GRANDPARENT_H
#define GRANDPARENT_H
#include "child.h"
...
#endif /* GRANDPARENT_H */
```

```
#define PI 3.14159
#define RADTODEG(x) ((x) * 57.29578)
#undef PI
```

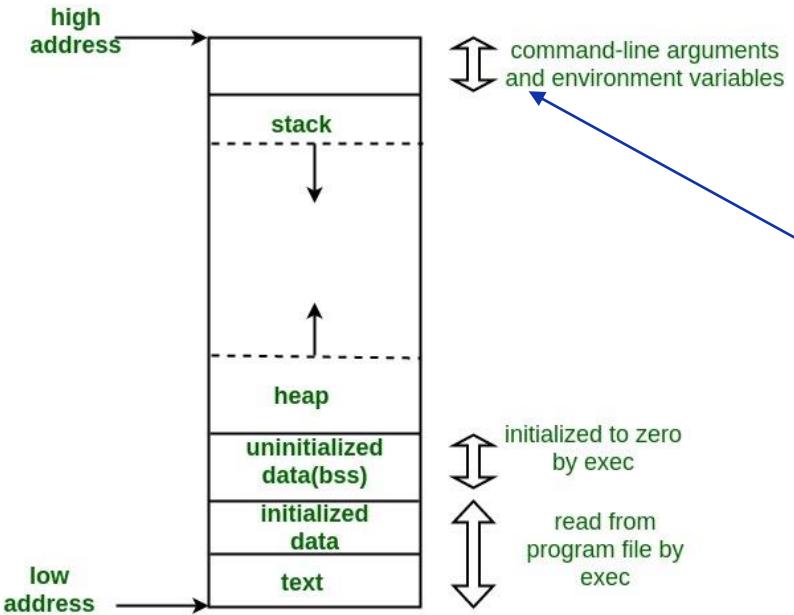
```
#if VERBOSE >= 2
printf("lots of trace messages\n");
#elif VERBOSE >1
printf("some trace messages\n");
#endif
```

```
#ifdef DEBUG
some_debug_function();
#endif
```

# Lecture 8 – Memory model and management

- Review
- Memory model of a C program
- Memory management

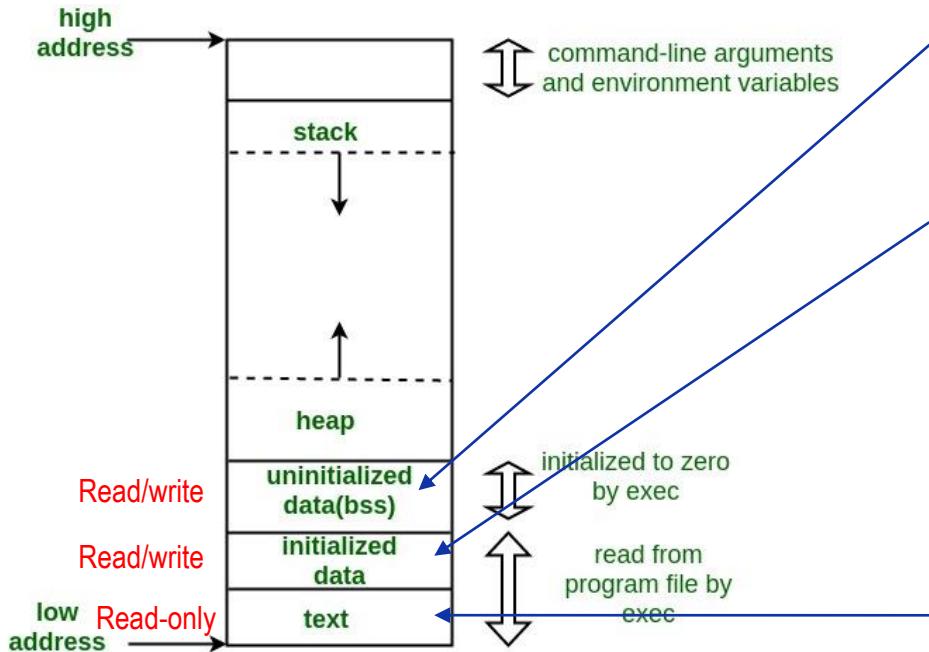
# Memory layout of C programs



- When a C program is running under an OS, it has its own memory space
  - OS takes care that the program does not exceed its memory boundaries
- In (bare metal) embedded systems, with no OS, the situation is similar, except
  - the segment for command line arguments and environment variables is not present

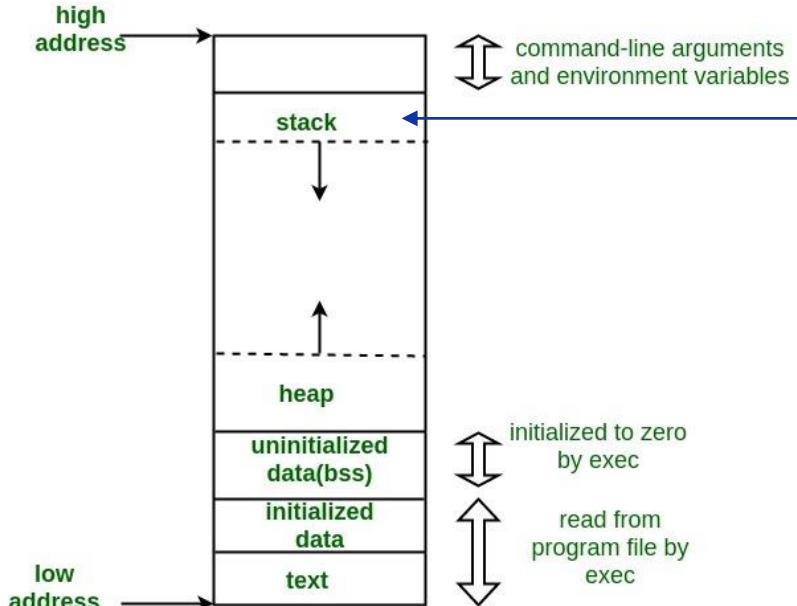
<https://www.geeksforgeeks.org/memory-layout-of-c-program/>

# Memory layout of C programs – static regions



- Uninitialized data segment (bss)
  - Contains global and static variables which are not initialized, or are initialized to zero
- Initialized data segment contains **initialized static variables**, that is, global variables and static local variables which have a predefined value
  - The size of this segment is determined by the size of the values in the program's source code, and does not change at run time
- Your program code is located in the "text" area of the memory
  - Read-only segment

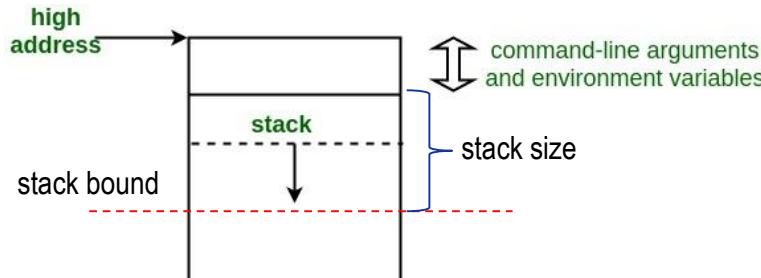
# Memory layout of C programs - stack



## What is a stack?

- At each function call, your program stores the values of local variables of a calling function to stack (often referred as "stack push"), after which the function is executed. When the function returns, the old values are "popped" from the stack (LIFO structure)
- Stack grows downwards
- The set of variables pushed/popped to/from stack is called a "stack frame"

# Memory layout of C programs – stack(2)

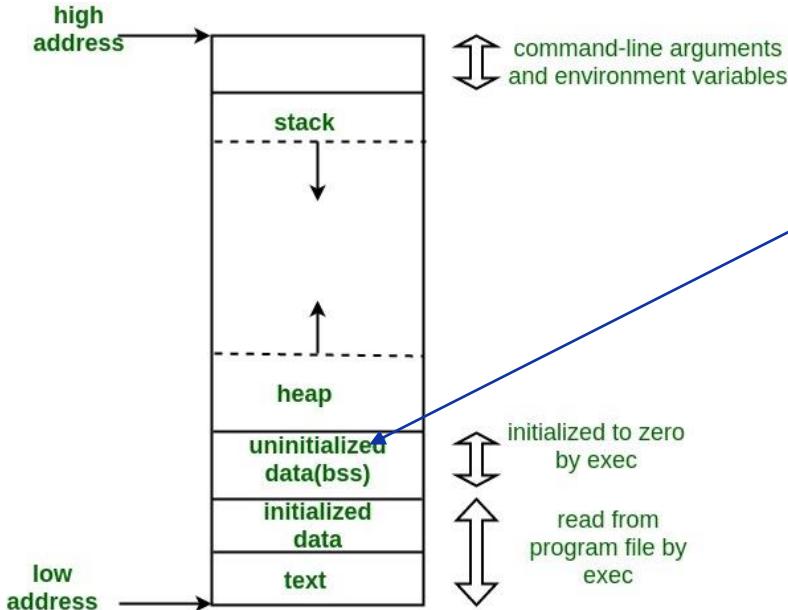


```
27  
28     IF :DEF: __STARTUP_CONFIG  
29 #include "startup_config.h"  
30     ENDIF  
31  
32     IF :DEF: __STARTUP_CONFIG  
33 Stack_Size EQU __STARTUP_CONFIG_STACK_SIZE  
34     ELIF :DEF: __STACK_SIZE  
35 Stack_Size EQU __STACK_SIZE  
36     ELSE  
37 Stack_Size EQU 2048  
38     ENDIF|  
39  
40     AREA STACK, NOINIT, READWRITE, ALIGN=3  
41 Stack_Mem SPACE Stack_Size  
42 __initial_sp
```

What is a stack overflow?

- A fatal error which occurs when the call stack pointer exceeds the stack bound. It is essentially a buffer overflow which causes your program to crash (A "segmentation fault")
- Typically caused by
  - infinite or very deep recursion
  - very large local variables
- Stack size is set at the beginning of your program – in the routines that are called before main(), often in assembler

# Memory layout of C programs - heap



What is a heap?

Like stack, heap is a fixed-size memory area, from where you allocate memory dynamically with malloc()  
Grows upwards

What is "out of heap space?"

Your system has ran out of memory – when the heap pointer has hit the stack pointer (or the heap boundary, if set)  
Often caused by badly behaving processes/threads which allocate memory, but do not free it after use (memory leakage)

# Lecture 8 – Memory model and management

- Review
- Memory model of a C program
- Memory management

# Dynamic memory allocation

- Two ways to allocate memory dynamically:
  - From **stack**
    - Named variables in functions
    - Allocated for you when you call a function
    - Deallocated for you when function returns
  - From **heap**
    - "Memory on demand"
    - You are responsible for all allocation and deallocation
  - The size of an array must be declared/defined before it can be used
  - Hence, the array may be insufficient or more than required to hold the data
  - To solve this issue, memory can be allocated dynamically from the heap
- ```
int x[10];
```
- ```
int* x = malloc(10 * sizeof(int));
```

# Using heap memory

- There are four functions, defined in `<stdlib.h>` for dynamic memory operations

<u>malloc()</u>	Allocates requested size of bytes and returns a pointer first byte of allocated space
<u>calloc()</u>	Allocates space for an array of elements, initializes to zero and then returns a pointer to memory
<u>free()</u>	deallocate the previously allocated space
<u>realloc()</u>	Change the size of previously allocated space

# malloc()

- The name malloc stands for "memory allocation"
- Reserves a block of memory of specified size and return a pointer of type void which can be casted into pointer of any form.
- If the space is insufficient, allocation fails and returns NULL pointer
  - This is "out of heap space" = "out of memory"
  - But, not a segmentation fault
- Warning: malloc() takes time, so be very careful in applications with hard real-time requirements!

```
ptr = (cast-type*) malloc(byte-size);  
  
ptr = (int*) malloc(100 * sizeof(int));
```

This statement will allocate either 200 or 400 bytes, depending on size of int (2 or 4 bytes), respectively. The pointer points to the address of first byte of memory

In C, auto-casting is possible, in C++ casting "(int\*)" to integer type is required

# calloc()

- The name calloc stands for "contiguous allocation"
- calloc() allocates multiple blocks of memory each of same size and sets all bytes to zero
- As with malloc(), if the space is insufficient, allocation fails and returns NULL pointer.

```
ptr = (cast-type*) calloc(n, element-size);  
  
ptr = (float*) calloc(25, sizeof(float));
```

This statement allocates contiguous space in memory for an array of 25 elements each of size of float, i.e, 4 bytes

In C, auto-casting is possible, in C++ casting "(int\*)" to integer type is required

# free(), realloc()

- Dynamically allocated memory created with either calloc() or malloc() doesn't get freed on its own
- You must explicitly use **free()** to release the space
- If the previously allocated memory is insufficient or more than required, you can change the previously allocated memory size using **realloc()**

```
free(ptr);
```

This statement frees the space allocated in the memory pointed by ptr

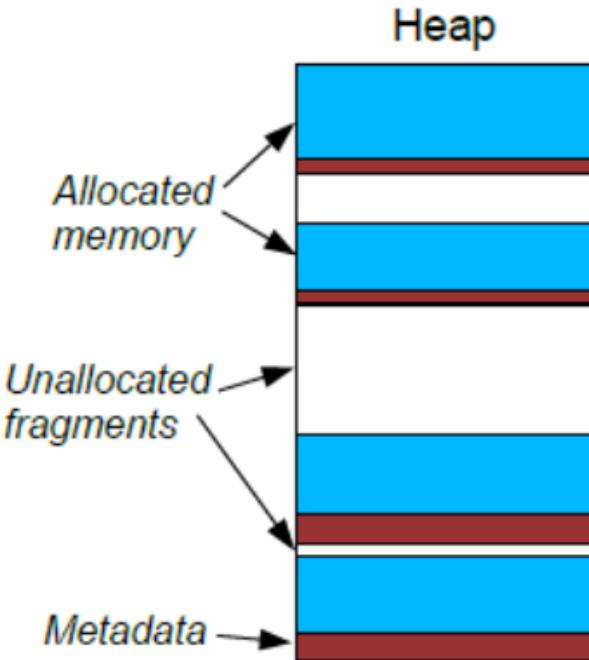
NOTE! After this you still have the pointer variable "ptr", but now it is a "**dangling pointer**". Attempt to dereference it (with \*) is a typical, fatal programming mistake.

```
ptr = realloc(ptr, newsize);
```

Here, ptr is reallocated with size of newsize

# A deeper look to malloc()

- You will actually reserve more memory than requested, because
  - Heap memory blocks may have minimum block size
  - Alignment requirements (padding)
  - Overhead of maintaining heap data structures
- malloc() includes an "allocator"
  - It keeps track where free memory is located – this is usually implemented using linked lists
  - The linked lists are "metadata", which is stored in heap as well, wasting heap memory space
  - Allocator tries to find the optimal memory block – the faster the allocator is, the sloppier work it does
- After heavy malloc() – free() – realloc() usage, there will be multiple chunks of heap memory that are unallocated. This is called fragmentation



# Useful memory manipulation functions

- Standard library <string.h> contains a set of useful functions for memory operations

```
void *memcpy(void *str1, const void *str2, size_t n)
```

Copies **n** characters from memory area **str2** to memory area **str1**

```
void *memmove(void *str1, const void *str2, size_t n)
```

Copies **n** characters from **str2** to **str1**, but for overlapping memory blocks, memmove() is a safer approach than memcpy()

```
void *memset(void *str, int c, size_t n)
```

Copies the character **c** (an unsigned char) to the first **n** characters of the string pointed to, by the argument **str**

# A deeper look into stack operation

- At the beginning of the program, a stack frame for main() is created
  - OS does this, or in case of no OS, start-up routines

```
int b() {  
    /* ... */  
}
```

```
int a() {  
    /* ... */  
    b();  
}
```



```
int main() {  
    /* ... */  
    a();  
}
```

Stack:

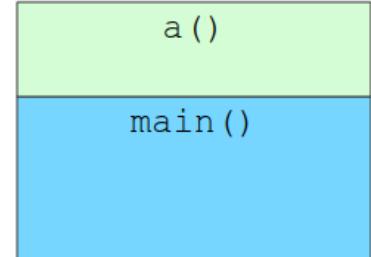
```
main()
```

# A deeper look into stack operation

- When a() is called, a new stack frame is created for a()

```
int b() {  
    /* ... */  
}  
  
int a() {  
    /* ... */  
    b();  
}  
  
int main() {  
    /* ... */  
    a();  
}
```

Stack:

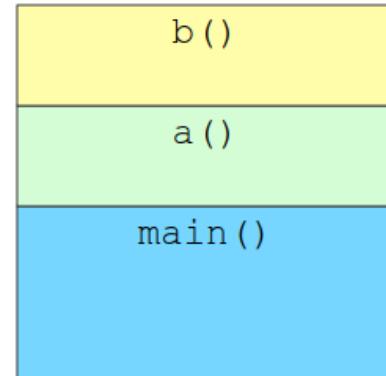


# A deeper look into stack operation

- Same for b: when b() is called, a new stack frame is created for b()

```
int b() {  
    /* ... */  
}  
  
int a() {  
    /* ... */  
    b();  
}  
  
int main() {  
    /* ... */  
    a();  
}
```

Stack:



# A deeper look into stack operation

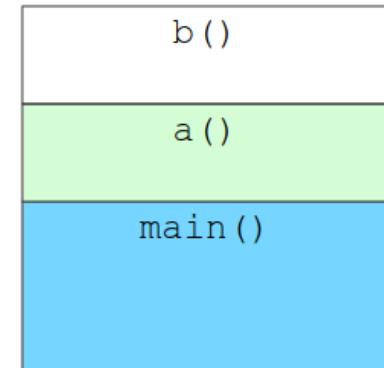
- When b() finishes running, its stack frame is removed
- ... and same happens when a() is finished

The diagram illustrates the state of the stack across three stages of function execution:

- Initial State:** The stack contains the main function frame (blue). A yellow arrow points from the code to this stage.
- After a() call:** The stack now contains both the main function frame (blue) and the a() function frame (green). A green arrow points from the code to this stage.
- After b() call:** The stack contains only the main function frame (blue). A dark blue arrow points from the code to this stage.

```
int b() {  
    /* ... */  
}  
  
int a() {  
    /* ... */  
    b();  
}  
  
int main() {  
    /* ... */  
    a();  
}
```

Stack:



# The map file

Image component sizes

- A report file, which shows
  - Memory segmentation
  - the data and code size for each module
  - And much more!

	Code (inc. data)	RO Data	RW Data	ZI Data	Debug	Object Name
	20	0	8	0	1304	adc.o
	2324	274	0	34	31599	app_timer.o
	160	22	0	12	6522	app_util_platform.o
	56	18	192	0	4096	arm_startup_nrf51.o
	1902	74	0	2	18044	diskio.o
	392	104	0	4	82783	error_handler.o
	2448	462	0	13	115595	main.o
	988	136	0	4	12700	sdcard.o
	1810	346	53	0	30380	spi_master_multislave_fast.o
	348	38	0	4	2287	system_nrf51.o
	972	226	28	0	7225	timers_init.o
	.....					
	66208	6378	1592	472	20904	Object Totals
	0	0	32	0	0	(incl. Generated)
	122	0	6	36	2	(incl. Padding)

m_gc	0x20006354	Data	16	fds.o(.bss)	Total RO Size (Code + RO Data)	72980 ( 71.27kB)
desc	0x20006364	Data	12	fds.o(.bss)	Total RW Size (RW Data + ZI Data)	21432 ( 20.93kB)
.bss	0x20006370	Section	44	localtime_w.o(.bss)	Total ROM Size (Code + RO Data + RW Data)	73040 ( 71.33kB)
_tms	0x20006370	Data	44	localtime_w.o(.bss)		
HEAP	0x200063a0	Section	2048	arm_startup_nrf51.o(HEAP)		
STACK	0x20006ba0	Section	2048	arm_startup_nrf51.o(STACK)		